

Server-Side Request Forgery (SSRF) — Part 3: Other advanced techniques



Maxime Leblanc [Follow](#)

Sep 11, 2017 · 5 min read

Hello geeks; This is the third and final part of this series about SSRF attacks. Let's remember that in [Part 1](#), we looked at the basic concepts behind such kinds of attacks. Then, in [Part 2](#) we learned how to circumvent protections that a website could put in place by using some IPv4 magic tricks from [Nicolas Grégoire](#). In this final part, I will explore with you two techniques I learned from [Orange Tsai](#) at the [DEFCON 25](#). All the credit for the present technique goes to him and I of course thank him very much for sharing his work with the community 😊. I also remind to you that all the application code I will be using for the following demonstrations is publicly available on my [SSRF Lab](#) hosted on GitHub and this last part is still using a [Docker-Compose](#) script that should get all of it up and running as-is within it's own network.

So, let's use one more time our Super Popular App presented in [Part 1](#) and [Part 2](#), which has only one purpose: It's a test page for WebHook integration, and should be able only to call external network addresses:



Also note that this time, our [WebHook code](#) is written in [Python Flask](#). We still have our

other `secret.corp` internal server located at `10.0.0.3`. Remember my conclusion of [Part 2](#), that could be rephrased as:

Do not try to parse URLs yourself; Instead use libraries that are made for this.

So you do exactly that, and using the usual [Python 2.7 libraries](#) you implement a filter on the `hostname` part of the url that was submitted:

```
url=request.form['handler']
host = urlparse.urlparse(url).hostname
if host == 'secret.corp':
    return 'Restricted Area!'
else:
    return urllib.urlopen(url).read()
```

Let's try to access this famous server by trying `http://secret.corp`:



SEE THE RESULT!

Restricted Area!

Well, the filter seems to do its job! Now, just for science, let's try `http://google.com#@secret.corp` instead (notice the whitespace character in the middle of the URL):

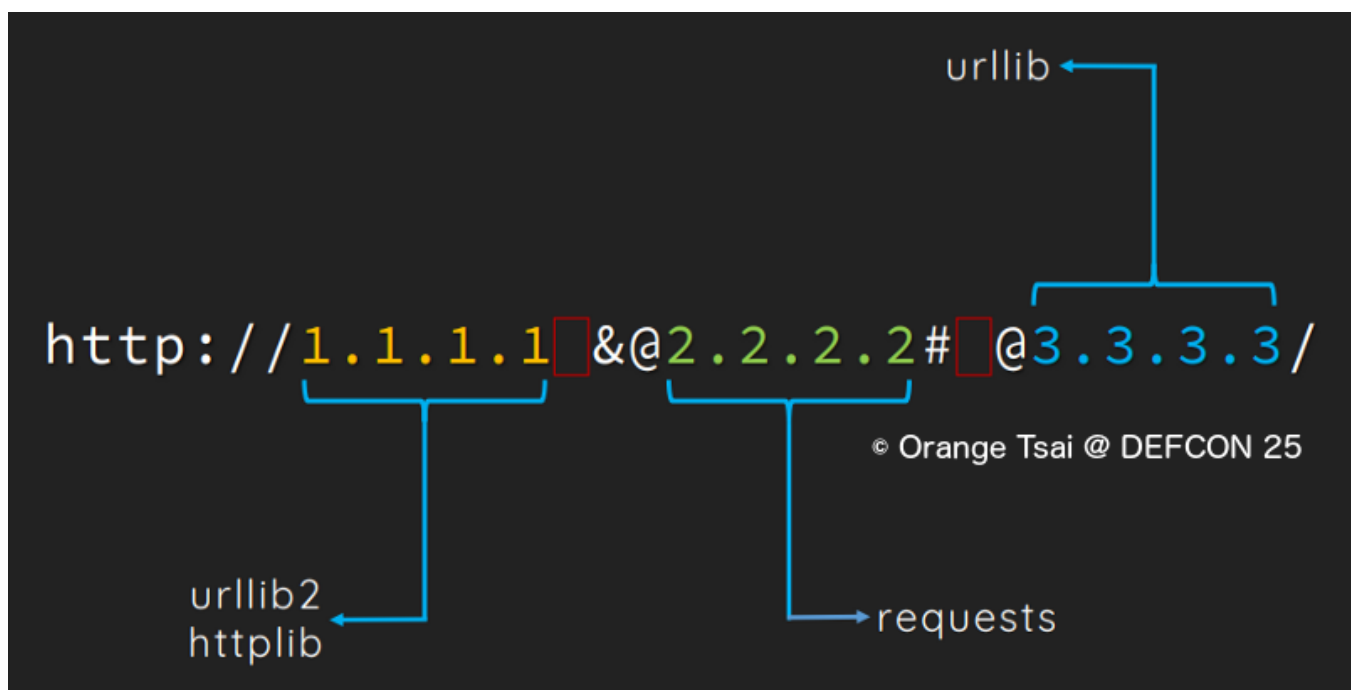


SEE THE RESULT!

This is secretserver!!
Only accessible from 10.0.0.0/8!

Wait, what happened again? Our filter using a standard parsing library got bypassed and authorized a request it wasn't supposed to? Well, yes and no; The problem here is that we use one library to filter the `hostname` and another one to do the actual request, and they do not interpret whitespaces in URLs the same way! The `urlparse` library recognized the URL as being addressed to `google.com` while `urllib` thought it was for `secret.corp` and acted accordingly. So here it's the difference of interpretation between the two libraries that lead to this SSRF vulnerability.

But it gets worse... This is part of a slide from Orange Tsai's presentation:



Source: Orange Tsai @ DEFCON 25

As you can see, Python libraries all have their own ways of handling whitespaces in URLs, and if you don't stick with one way throughout your whole application, unexpected results might happen when dealing with those. Here is what says the official RFC3968:

The authority component is preceded by a double slash (“//”) and is terminated by the next slash (“/”), question mark (“?”), or number sign (“#”) character, or by the end of the URI

From this definition, it is kind-of ambiguous which library behaves right; I would personally say that none correspond to this behaviour exactly.

Future work: Protocol injection

As a conclusion, I would like to introduce the notion of protocol injection: Remember in [Part 1](#) the structure of a URL:

scheme://user:pass@host:port/path?query=value#fragment

How could-it be possible to bypass a fixed `scheme://` with only access to the HTTP(S) protocol? Meet the newline injection technique. Now, at the time of this writing, I have not implemented it as a webserver exploit in my SSRF Lab, but it will be available on this [GitHub repository folder](#) in a short time. Meanwhile, let's consider a user on our WebHook testing server. A mail server is located at `mail.corp`, using a default Ubuntu Postfix installation. Our WebHook server can only make HTTP requests. However, if we are able to inject newline characters (`\r\n`) as in the hostname and the library accepts it, it is possible to send valid commands to the SMTP server during the SSL handshake. This is what the TCP response from the mail server looks like:

Request:

url = https://mail.corp\r\nHELO web.corp\r\nMAIL FROM...:25/

Response:

SMTP: 502 5.5.2 Error: command not recognized <SSL Gibberish>
Computer Security Hacking Domain Names

SMTP: 250 2.1.0 Ok

...

SMTP: 502 5.5.2 Error: command not recognized <SSL Gibberish>

About Help Legal

The trick here is that during the SSL handshake, the hostname is used entirely, including its newlines as “end-of-command” signals so we can insert valid commands in between the SSL handshake traffic. I still have some testing to do in order to have a valid POC of this in my [Lab](#), but my first [Wireshark](#) captures are promising 😊. Note here that the same trick should apply for any text-based protocol : I can imagine well using this in order to register to an internal SIP server, for example. Leave a comment if you have a suggestion of implementation for me and final thanks to [Orange Tsai](#) and [Nicolas Grégoire](#) for their work!

Happy Hacking!

